



Presentation Slide Deck

AI Fundamentals Glossary

90+ AI terms explained in two minutes each



35 lessons · Beginner



WHAT'S INSIDE

35 lessons


1h 31m

Beginner



- 01 **Fundamentals**
AI, machine learning, GPUs, language
- 02 **Predictive AI**
AutoML, forecasting, foundation models
- 03 **Generative AI**
LLMs, tokens, RAG, prompting
- 04 **Agentic AI**
Agents, orchestration, MCP, context
- 05 **Observability**
MLOps, monitoring, drift, evaluation
- 06 **Governance**
Data, AI and agentic governance
- 07 **Platform & Compliance**
Deployment models, four regions

One arc, seven stages



Every stage has its own vocabulary. The videos follow this order, and so does this deck.

Compliance sits across all of it. The last four lessons cover the United States, Europe, Asia-Pacific and Latin America.



What everything else is built on

Artificial Intelligence

Systems that learn from data instead of following fixed rules

Machine Learning

Learning to tell cases apart from thousands of examples

Features

The variables that tell a model what to pay attention to

GPU

Parallel compute that makes training on text and images viable

Neural Networks

Layers of nodes that adjust connections from their own errors

Deep Learning

Stacked layers learning steadily more abstract representations

Transformer

Attention: weighing every part of an input against every other

NLP and NLU

Processing language, then grasping intent and context

AGI

A hypothetical AI that reasons across any domain a human can

Turning history into decisions

Predictive AI

Acting before events happen rather than reporting after

AutoML

Automated algorithm selection and tuning, hours instead of weeks

Tabular Foundation Model

Pre-trained on structured data, so no per-task training

Supervised Learning

Training on examples where the correct answer is known

Unsupervised Learning

Finding natural groupings with no predefined categories

Classification

Predicting which category something belongs to

Regression

Predicting a continuous value rather than a category

Clustering

Grouping by similarity, discovering rather than classifying

Deep Learning at Scale

Where depth and data beat hand-built rules

From raw data to a live prediction

Forecasting

Historical data projected forward across many variables

Time Series Models

Built for sequential data, so seasonal patterns stay visible

Anomaly Detection

Flagging when reality diverges from the expected pattern

NLP Prediction

Deriving a label or score from text rather than generating it

Vision Modeling

The same principle applied to images and video

Feature Engineering

Reshaping raw data into what a model can actually learn from

Data Science Lifecycle

Problem framing through to monitoring, every stage connected

Model Scoring

Inference: predictions on data the model has never seen

Where ML Still Wins

Regulated, high-volume and bespoke problems

Models, and what they cost to run

Generative AI

AI that creates text, code and answers rather than only analysing

GPT

Generative Pre-trained Transformer, the architecture behind most LLMs

Proprietary LLMs

Fast to start, priced per token, data leaves your environment

Open Source LLMs

Open weights, self-hosted, no per-token cost, you run the stack

Token

The chunk of text a model processes, not a word or a character

Tokenomics

The economics of token usage, and why cost compounds at scale

NLG

Producing human-readable text from data or instructions

NER

Pulling people, dates and values out of unstructured text

SLM

Fewer parameters, trading breadth for precision and speed

Retrieval, grounding and quality

Information Retrieval

Finding relevant content from a large collection

RAG

Retrieve at question time, then answer from what was retrieved

Graph RAG

Retrieval through a knowledge graph, so relationships are navigable

Agentic RAG

Retrieval as active reasoning: rewrite, route, evaluate, retry

Hallucination

Plausible text filling a gap the model could not actually answer

Grounded Generation

Answers anchored in verified sources, so they stay traceable

LLM Evaluation

Accuracy, relevance, coherence, cost and safety, measured over time

RAG Evaluation

Retrieval and generation assessed separately, because either can fail

Tracing Matrix

Mapping inputs to outputs to isolate where quality breaks

Steering models, and putting them to work

Prompt Engineering

Designing inputs that produce useful, well-structured output

System Prompt

Instructions set before the conversation, governing everything after

Prompt Catalog

Tested, reusable prompts shared across teams

Fine-Tuning Engine

Training a pre-trained model further on your own domain data

LLMOps

Deploying, monitoring and updating language models over time

Conversational AI

Multi-turn dialogue that holds context and escalates cleanly

Document Extraction

Structured data pulled out of contracts, claims and filings

Text-to-SQL

A plain English question turned into a database query

SLM Development

Building a compact model for one well-defined task

Systems that act on a goal

AI Agent

Given a goal, it works out the steps itself and carries them out

Agentic AI

Operating across many steps, deciding and adapting along the way

Orchestration

The logic deciding which component runs, in what order

Agent Orchestration

Autonomous components handing work to each other mid-task

Agent Runtime

The environment where agents actually execute

Agent Harness

Agents, tools, memory and instructions held in one framework

Deep Research Agent

Searches, reads and cross-references into a structured answer

Coding Agent

Writes, tests and debugs its own code until it works

Vertical Agents

Built for one industry, which is what makes them precise

How an agent reaches the world

Plan



Act



Check



Self-correct



Deliver

The internal loop that separates an agent from a chatbot

Agent Tools

The actions available: search, code, documents, APIs, audio

Tool Builder

Custom tools reaching internal databases and legacy systems

Agent Builder

Where the goal, tools, instructions and limits get configured

MCP

One standard for connecting models to external tools and data

MCP Runner

The runtime managing those connections reliably at scale

Knowledge Graph

Entities and relationships an AI system can reason across

Ontology

The definitions and categories giving the graph its logic

General-Purpose Agent

Chooses research, code or retrieval as the task demands

Keeping it working after launch

MLOps

Deployment, versioning, monitoring and retraining as one practice

LLMOps

The same, plus prompts, token cost and context window limits

Model Hub

One place models are published, versioned and deployed from

Observability

Understanding a system from what it produces, by design

Monitoring

Tracking defined metrics and alerting on a crossed threshold

Model Monitoring

Accuracy, confidence, output distribution, latency

Agent Monitoring

Which actions agents take, and whether they stay in bounds

Telemetry

The logs, traces and metrics everything above depends on

Evaluation

Measuring quality continuously, not once before launch

Staying accountable at scale

Drift

Production data drifting away from what the model was trained on

Model Drift

The resulting decay, dangerous because it is quiet

Risk Management

Bias, security, regulatory exposure and edge-case behaviour

ML Interpretability

Showing why a decision was made, so it can be challenged

Data Governance

Who owns data, who can reach it, how long it is kept

AI Governance

How models get approved, audited and answered for

Agentic Governance

Boundaries, audit logs and human oversight for autonomous agents

Responsible AI

Fair, transparent and accountable across the whole lifecycle

Agent Evaluation

Judging decisions and actions, not just the text produced

Where it runs, and the rules that apply

On-Premises

Inside your own data centres,
no data leaving the building

Air-Gapped

Cut off from every external
network, including the internet

Cloud, SaaS, VPC

Scale on shared infrastructure,
or a private slice of it

Sovereign AI

Data, models and compute held
inside one jurisdiction

Kubernetes

Automating deployment and
scaling wherever the workload
runs

Systems of Record

Where authoritative structured
data lives

System of Intelligence

The layer that makes that data
useful

Low-Code AI

Removing the barrier that kept
AI out of most of a business

COMPLIANCE, FOUR REGIONS

United States

SOC 2 · HIPAA · FedRAMP High

Europe

GDPR · EU AI Act · ISO 27001

Asia-Pacific

PDPA · PIPL · Privacy Act · DPDP

Latin America

LGPD · Law 1581 · LFPDPPP

90+ terms. Seven parts. One vocabulary.

■ Take the quiz

Earn your H2O.ai University certification

■ Go deeper

Full courses and learning paths at h2o.ai/university

■ Keep this deck

One page per part, for the term you cannot place

h2o.ai/university

